SingularityNeo — Stakeholder & Sales Pitch

> *Dual-purpose deck: use as internal leadership alignment and as an external sales narrative.*
> *Each section is one slide. Bullets are slide content. Italic blocks are speaker notes.*
> *Recommended format: 16:9, 16 slides, 20-minute read-through.*

Slide 1 — Title

SingularityNeo

The AI-Native Cockpit for Engineering Teams

Provable autonomy. Every agent action carries its receipts.

Speaker note: Open with the tagline. Don't describe features. Set the frame that this is about trust + autonomy, not another copilot.

Slide 2 — The Moment

AI is about to hit the autonomy trust cliff.

- Agents can now write code, run tests, open PRs, transition tickets
- But engineering leaders cannot answer four questions:

1. **What was the agent allowed to do?**
2. **What context did it actually read?**
3. **What decision did it make, and why?**
4. **Can I replay it for my auditor, my CISO, my board?**

- Without those answers, autonomy stalls at "toy project" scale

Speaker note: Every buyer in this space is nervous. Naming the fear gives us permission to sell the cure.

Slide 3 — The Market Has Three Camps — None of Them Solve It

Camp	Examples	Optimized For	What's Missing
Coder Copilots	Cursor, Copilot, Windsurf	Keystroke speed	No org intent, no memory, no audit
Autonomous Agents	Devin, Factory, Replit Agent	Task completion	Opaque reasoning, "trust me"
Work Trackers + AI	Linear, Jira, ServiceNow AI	Governance, tracking	Agents can't do the work

No one fuses execution + governance + evidence in one substrate.

Speaker note: Draw the triangle on the whiteboard if live. Pause. Let the gap land.

Slide 4 — The Differentiator

Replayable Provenance,

bound to a Signed Capability Contract.

Every agent action in SingularityNeo is:

1. **Bound** to a human-signed capability contract
2. **Executed** under a declared role policy (tools, memory, quality bar)
3. **Written** to a unified, append-only timeline
4. **Packaged** as a content-addressed, replayable evidence bundle

Not a compliance layer. The substrate.

Speaker note: This is the core slide. Memorize it. Say it the same way every time.

Slide 5 — What That Means in One Sentence

> *"Yes, agents shipped 40% of this quarter's work — and here is the exact contract they worked to, the exact tools they were allowed to call, the exact context they read, the exact decisions they made, and I can replay any of it in one click."*

That sentence is what engineering leaders want to say. SingularityNeo is the only product that lets them say it truthfully.

Speaker note: Ask the room: "Can you say this today about your AI tools?" Silence sells the product.

Slide 6 — Four Substrate Ingredients. One Brain.

Ingredient	What It Is	Why It's Hard
Capability Contract	Human-signed outcome, DoD, boundary, owner — before any run	Requires hierarchy + versioning, not a prompt
Role Policy	Declared tool allow/deny + memory scope + quality bar per agent role	Requires runtime enforcement, not prompt tricks
Unified Timeline	Chat, thoughts, tool calls, retrievals, artifacts, approvals — one stream	Requires fusing runtime telemetry into the UX
Evidence Packet	Content-addressed, replayable bundle, shareable via URL	Requires deterministic capture from day one

Cursor has #3 partially. Devin has #2 partially. Jira has #1 partially. Nobody has all four.

Speaker note: The "partially" column is the punchline. Competitors can't bolt these on — they'd have to rewrite their core.

Slide 7 — The Product Story: One Cockpit

Not a dashboard. Not a chat sidebar. **A cockpit.**

- One work item = one session
- Plan, chat, runs, tools, evidence, learning — **one frame, one stream**
- Every agent turn shows: role · tools used · context read · decision · citations
- A run is not a different page — it's a live block in the timeline
- No tab-switching to approve, upload, or read an artifact

The competition is "many tabs." We are one brain.

Speaker note: If doing a live demo, this is where you open the Cockpit. Don't narrate features — just use it. Let the unified stream sell itself.

Slide 8 — The Ten-Minute Wow

Golden path from install to shareable evidence

1. Install / open app — **30 s**
2. Connect GitHub repo — **60 s**
3. Approve workspace root — **30 s**
4. Agents read repo, propose capability contract — **90 s**
5. Pull a real issue as a work item — **60 s**
6. Run the default workflow — **3 min**
7. **Evidence packet generated, permalink copied to clipboard**

Paste the permalink into a PR or Slack. That link* is *the marketing.

Speaker note: Prospects never forget the moment the evidence URL lands in Slack. Make them experience it.

Slide 9 — Agent Differentiation You Can See

Every built-in role has **declarative, enforced** differences:

- **OWNER** — approves contracts, no direct code writes
- **ARCHITECT** — writes to capability memory, full read scope
- **BUILDER** — code + test tools, cannot self-approve
- **REVIEWER** — read-only, must cite ≥3 sources, cannot modify files
- **CRITIC** — contrarian mode required, must disagree before signing off
- **OPERATOR** — runtime only, no design tools

Role chips appear on every agent turn in the timeline. Difference is observable, not implied.

Speaker note: "Why do I need five agents instead of one ChatGPT?" This is the answer. They're not different prompts — they're different policies.

Slide 10 — Safety That Precedes Orchestration

The Capability Readiness Contract

A workflow cannot run until six gates are green:

```
✓ Repo connected
✓ Workspace root approved
✓ Owner assigned (human)
✓ Outcome statement
✓ Definition of Done (≥3 testable checks)
✓ Minimal contract published
```

Below 6: agents can reason, summarize, critique — but cannot act. At 6: autonomy unlocks, with policy enforcement and evidence capture.

Speaker note: This is why our agents don't go rogue. The gate is the product.

Slide 11 — The Evidence Packet

Every run produces a shareable, replayable bundle

Contains:

- The capability contract the run executed against
- The role policies that were enforced
- Every tool call, with inputs and outputs
- Every retrieval, with sources
- Every decision, with citations
- Every artifact, content-addressed
- Human approvals and overrides

Exportable to ALM, auditors, regulators. Permalink for PRs and Slack.

Speaker note: Hold up a printed evidence packet if in person. Physical object sells digital provenance.

Slide 12 — Why Now

- **Model capability is commoditizing** — autonomy is no longer science fiction
- **Regulation is catching up** — EU AI Act, SR-11-7, ISO 42001 all require decision traceability

- **Enterprise buyers are burned** — 2024–2025 pilots exposed the trust gap
- **Coding tools are table stakes** — the next buying cycle is about governed autonomy
- **The substrate takes 12–24 months to build** — the window to own this is now

Speaker note: If asked "why won't OpenAI/Cursor do this?" — they can't, without rewriting. The substrate is the moat.

Slide 13 — The Moat

It isn't the AI. It's the substrate.

- Capability hierarchy
- Role policy runtime
- Unified interaction timeline
- Content-addressed evidence store
- Readiness contract as a first-class object

Competitors started with a chat box or a code editor. We started with the ledger.

You can add a chat box to a ledger. You can't add a ledger to a chat box.

Speaker note: This is the line to close an investor or a skeptical CTO. Deliver it slowly.

Slide 14 — Who Buys

Primary buyer: VP Engineering / Head of Platform

They feel three pains:

1. Scattered AI tool sprawl with no audit trail
2. CISO and Risk blocking broader AI rollout
3. Pressure to show AI-driven velocity gains to the board

We give them:

- One cockpit that consolidates the sprawl
- An evidence trail that unblocks CISO
- A quarterly "Agent Review Report" that ships to the board

Secondary champions: Engineering Managers, Staff Engineers, Compliance

Speaker note: Sell to the VP. Demo to the Staff Engineer. Ship proof to the CISO.

Slide 15 — The Business Case

Metric	Without SingularityNeo	With SingularityNeo
AI adoption ceiling	Pilot / sandbox	Production-wide
Audit readiness	Manual reconstruction	One-click export
Agent-shipped work	Unreported / unsafe	Measured + governed
Tool sprawl	4–7 tools	1 cockpit
Time-to-first-value	Weeks	**10 minutes**

The ROI isn't "faster coding." It's "AI you can actually deploy org-wide."

Speaker note: Don't compete on lines-of-code savings. Compete on the autonomy ceiling you just raised.

Slide 16 — Call to Action

See it in 10 minutes.

- **Prospects:** Live golden-path demo — your repo, your issue, a real evidence packet
- **Stakeholders:** 2-week pilot with one team, one capability, measurable before/after
- **Investors / Board:** Quarterly Agent Review Report from the pilot team

One line to leave the room with:

> ***Everyone else sells agent speed. We sell agent receipts — and in enterprise engineering, receipts are what actually unlock adoption at scale.***

Speaker note: End on the receipts line. Don't add anything after it.

Appendix A — One-Slide Version (if you only get 60 seconds)

> *SingularityNeo is the AI-native cockpit for engineering teams.*
> *We're the only product where every agent action is bound to a human-signed capability contract, executed under an enforced role policy, captured in a unified timeline, and packaged as a replayable evidence bundle.*
> *Cursor sells speed. Devin sells autonomy. Jira sells governance.*
> *We sell **provable autonomy** — and that's what unlocks AI adoption at enterprise scale.*

Appendix B — Objection Handling

Objection	Response
"We already have Copilot."	Copilot is a keystroke tool. It has no contract, no role policy, no evidence trail. Use it *inside* the Cockpit.
"Can't you just add audit to Cursor?"	You can add logs. You can't add a capability hierarchy, role-enforced tool policy, or a signed contract model without rewriting.
"Our compliance team will never approve AI agents."	Start with the evidence packet. That's the artifact your compliance team has been asking every other AI vendor to produce.
"How is this different from LangSmith / Langfuse?"	Those are developer-facing traces. This is a product end users (engineers + leaders) live in. Traces are an ingredient; the cockpit is the meal.
"Why not build in-house?"	You can. It will take 18–24 months, two platform teams, and the loss of the market window.
"What about open-source frameworks (LangGraph, AutoGen)?"	Frameworks for building agents. Not products for running a team on agents. Different category.

Appendix C — Demo Script Skeleton (10 minutes)

1. **(0:00)** Open the app — show three surfaces only: Home, Cockpit, Studio
2. **(0:30)** Connect GitHub → approve workspace
3. **(1:30)** Agents propose capability contract from repo — approve it live
4. **(3:00)** Pull a real GitHub issue as a work item, land in the Cockpit
5. **(3:30)** Type one sentence of intent — workflow begins
6. **(4:00) Show the timeline:** role chips, tool calls expanding inline, knowledge lens updating on the right
7. **(6:00)** Reviewer agent disagrees with the Builder — show contrarian mode in action
8. **(7:00)** Human approves the decision in the timeline — no tab switch
9. **(8:00)** Run completes → evidence packet permalink auto-copied
0. **(8:30)** Paste permalink in a Slack channel → replay the entire run from the link
1. **(9:30)** Show Studio once to prove agent role table exists — close it
2. **(10:00)** "That's it. One cockpit. One brain. One receipt per run."

Document version: 1.0 · For internal stakeholder review and external sales conversations · Derived from product strategy sessions